

PARITY

An Adaptive Horror Game Driven by Online Player-Trajectory Prediction

Project Review III — Progress Report Shivam Kumar · Vellore Institute of Technology · 16 September 2026

1. Domain and Problem Statement

1.1 Domain

Game AI has largely optimised for *competence* — agents that play well. A separate and less-solved problem is **adaptation**: systems that model the individual player and reshape the experience around them. This sits at the intersection of three areas: player modelling (inferring player state and intent from behavioural telemetry), adaptive game AI and dynamic difficulty adjustment (acting on that inference), and procedural content generation (changing the world itself rather than only the agents in it).

Horror is the genre where this matters most, and where it is least explored. Horror depends on *uncertainty*. Every horror game degrades as the player learns it: routes get memorised, spawn points get learned, and the monster becomes a timing puzzle. The genre's central failure mode is that mastery destroys the product.

1.2 Problem Statement

Existing horror antagonists are **static or randomised**. Static AI is learnable and stops being frightening. Randomised AI is unlearnable and therefore reads as unfair rather than intelligent. Neither adapts to the specific player.

PARITY proposes a third option: an antagonist that builds an **online behavioural model of the individual player** and acts on its predictions — intercepting where the player is going rather than chasing where they are, and degrading the information the player has come to depend on.

1.3 Why this is a hard and interesting problem

1. **Cold start.** A commercial ML pipeline trains offline on a corpus. A horror antagonist must be interesting in the *first* two minutes, against a player it has never seen, with no pre-collected data for that individual.
2. **Legibility.** A model the player cannot perceive is, to the player, no model at all. The adaptation must be *felt* without being announced.
3. **Fairness.** An adaptive antagonist that is merely stronger is not scarier, it is unfair. The system needs explicit constraints that keep it defeatable.
4. **Evaluation.** "Is it scary?" is not measurable at the scale a solo project can playtest. The project needs proxy metrics that can be computed headlessly.

1.4 Aim and Objectives

Aim: Build a first-person horror game in which the antagonist's behaviour is driven by a live, per-player predictive model, and demonstrate quantitatively that the model both learns and *matters*.

#	OBJECTIVE	STATUS
O1	Formalise the facility as a mutable room graph with runtime topology edits	Done
O2	Implement an online next-room predictor that learns from the first move	Done
O3	Show it beats a no-model control and an order-1 Markov baseline	Done
O4	Show it learns <i>nothing</i> from a random player (no leakage)	Done
O5	Show better prediction yields a measurably more dangerous antagonist	Done
O6	Make the model legible to the player in-game	Demo built
O7	Ship the first-person Godot vertical slice	In progress

2. Literature Review

Twenty-two works, twenty of them from 2023 or later, grouped by the question they answer for this project.

2.1 Player modelling and behaviour prediction

#	WORK	YEAR	RELEVANCE TO PARITY
1	Lin et al., <i>Open Player Modeling: The Open Player Socially Analytical Intelligence Architecture</i> — arXiv:2603.26915	2026	Argues player models should be inspectable rather than hidden; directly motivates our in-game prediction overlay and death readout.
2	Dehpanah et al., <i>Player Modeling using Behavioral Signals in Competitive Online Games</i> — arXiv:2112.04379	2021	Establishes that raw behavioural signals alone support useful player models — no questionnaires or biometrics needed.
3	Carlier et al., <i>Personalization in Serious Games and Gamification for Healthcare</i> — arXiv:2411.18500	2024	Three-tier taxonomy of personalisation (model / method / adaptation) that our Telemetry–Predictor–Director split mirrors.
4	Halina et al., <i>Representing and Generating Levels Over Time through Playtrace Reconstructive Partitioning</i> — arXiv:2607.12097	2026	Uses playtraces as the primary representation of a level — the same premise as modelling the facility through the player's path.
5	Bazzaz, <i>Player Perceptions of Generative AI in Games: A Steam Review Analysis</i> — arXiv:2608.11539	2026	Evidence on how players actually react to AI-driven systems; informs how openly PARITY should surface its model.

2.2 Adaptive game AI and dynamic difficulty adjustment

#	WORK	YEAR	RELEVANCE TO PARITY
6	Romeo et al., <i>NTRL: Encounter Generation via Reinforcement Learning for DDA in Dungeons and Dragons</i> — arXiv:2506.19530	2025	Closest analogue: adapts <i>encounters</i> rather than agent stats. PARITY adapts encounter placement via prediction instead of RL.
7	Lopes et al., <i>Closing the Loop in Affect-Driven Game Adaptation: A Systematic Review</i> — arXiv:2505.01351	2025	Survey of the affect-adaptation loop; identifies evaluation as the field's weak point — the gap our headless bot harness addresses.
8	Cafri et al., <i>Dynamic Difficulty Adjustment With Brain Waves</i> — arXiv:2504.13965	2025	Represents the biometric branch of DDA. We deliberately avoid biometrics: behaviour alone, no extra hardware.
9	Fuchs et al., <i>Personalized Dynamic Difficulty Adjustment — Imitation Learning Meets Reinforcement Learning</i> — arXiv:2408.06818	2024	Learns to imitate an individual player, then adapts. Same personalisation goal; far heavier training requirement than ours.
10	McConnell et al., <i>From Frustration to Fun: An Adaptive Problem-Solving Puzzle Game Powered by Genetic Algorithm</i> — arXiv:2509.23796	2025	Evolutionary adaptation of content difficulty; a design-space alternative to our utility-scoring Director.
11	Jarma Montoya et al., <i>Atari Games Challenge: Multimodal Player Experience Assessment</i> — arXiv:2605.27261	2026	Multimodal experience measurement; the evaluation vocabulary for the later playtest milestone.
12	Tripathi et al., <i>AI-Enabled Serious Games: Integrating Intelligence and Adaptivity in Training Systems</i> — arXiv:2605.21962	2026	Architectural patterns for embedding adaptive intelligence in a game loop.

2.3 Procedural generation of mutable levels

#	WORK	YEAR	RELEVANCE TO PARITY
13	Fiuza Vieira et al., <i>A Novel Procedural Generation for Level Design of Mansions and Dungeons</i> — arXiv:2606.03857	2026	Graph-based generation of exactly our building type; informs the room-graph generator and connectivity guarantees.
14	Özkan, <i>Procedural Game Level Design with Deep Reinforcement Learning</i> — arXiv:2510.15120	2025	RL-driven level construction; contrasts with our constraint-checked mutation approach.
15	Earle et al., <i>Video Game Level Design as a Multi-Agent Reinforcement Learning Problem</i> — arXiv:2510.04862	2025	Treats level design as multi-agent decision-making — the framing behind our Director editing the level at runtime.

2.4 Sequence and next-location prediction (methodological core)

#	WORK	YEAR	RELEVANCE TO PARITY
16	Begleiter, El-Yaniv & Yona, <i>On Prediction Using Variable Order Markov Models</i> , JAIR 22:385–421 — arXiv:1107.0051	2004	Foundational. The formal treatment of VOMM prediction our model implements, including back-off/interpolation behaviour.
17	Wu et al., <i>Beyond Regularity: Modeling Chaotic Mobility Patterns for Next Location Prediction</i> — arXiv:2509.11713	2025	Directly analogous task — next-location prediction over a discrete place graph — and handles the irregular case our noisy player represents.
18	Deng et al., <i>STRelay: A Universal Spatio-Temporal Relaying Framework for Location Prediction</i> — arXiv:2508.16620	2025	State-of-the-art trajectory prediction architecture; the natural upgrade path from VOMM in the stretch milestone.
19	Liu et al., <i>Mixture-of-Experts for Personalized and Semantic-Aware Next Location Prediction</i> — arXiv:2505.24597	2025	Per-individual personalisation of location prediction — the same personalisation axis PARITY needs.
20	Soni et al., <i>Can We Predict Your Next Move Without Breaking Your Privacy?</i> — arXiv:2507.08843	2025	Privacy framing for behavioural trajectory models; relevant because PARITY persists a per-player model across sessions.

2.5 Fear and horror as a measurable experience

#	WORK	YEAR	RELEVANCE TO PARITY
21	Zhang et al., <i>VRMN-bD: A Multi-modal Natural Behavior Dataset of Immersive Human Fear Responses</i> — arXiv:2401.12133	2024	Evidence that fear responses are detectable in ordinary behavioural signals — supports using movement telemetry as a tension proxy.
22	Zhang et al., <i>Decoding Fear: Exploring User Experiences in Virtual Reality Horror Games</i> — arXiv:2312.15582	2023	Qualitative account of what actually frightens players; sources the design constraint that unfairness destroys fear.

2.6 Research Gap

Across these works, three gaps are consistent:

1. **DDA adapts parameters, not position.** The literature overwhelmingly tunes difficulty *magnitude* (enemy health, encounter budget, puzzle complexity). Almost none adapt *where the antagonist chooses to be* from a predictive model of the individual player.
2. **Next-location prediction is not used adversarially.** Trajectory prediction is mature (§2.4) but is applied to service delivery and urban analytics. Its use as a real-time adversary against the person being predicted is unexplored.
3. **Adaptation is invisible and therefore unevaluated.** Lopes et al. (2025) identify evaluation as the field's weak point. Adaptive systems are rarely made legible to the player or measured against a no-model control.

PARITY addresses all three: it applies online trajectory prediction adversarially, adapts antagonist *placement* and *world state* rather than parameters, and evaluates against explicit controls with the model surfaced on screen.

3. Proposed Methodology

3.1 Formalisation

The facility is an undirected graph $G = (V, E)$: rooms are vertices, doors are edges. The player occupies a vertex and moves along edges, producing a trajectory $h = (v_1, v_2, \dots, v_t)$.

The clock. The facility runs on a single shared cycle: lockdown, then a short window in which every door unlocks, then lockdown again. Decisions happen only at windows. During lockdown the player is sealed in place — which is when they draw in the notebook, copy a codebook, or stage the outgoing register. Drawing is real-time, so an open door and an unfinished page compete for the same seconds.

Prediction task. At each window the player chooses a door *or declines to move*. Given trajectory h , estimate

$$P(v_{t+1} = c \mid h) \text{ for each } c \in N(v_t) \cup \{v_t\}$$

Staying is a first-class action, not the absence of one. Including v_t in the candidate set is what makes dwell predictable at all — and dwell is highly informative, because a player lingers where there is something to record. The antagonist is bound by the same clock: it also acts only at windows, so lockdown is genuinely safe and all the danger concentrates at the moment the doors open.

Why a variable-order model is required. An order-1 Markov model conditions only on v_t . But real players run *routes*, so the same room has different successors depending on how it was entered. In our test facility, room D appears five times in a habitual player's circuit with four different successors — order-1 is structurally incapable of separating them. Conditioning on longer context resolves it, but long contexts are sparse. A **variable-order Markov model with back-off** resolves both pressures at once.

3.2 The model

For each context length $o = 0 \dots k$ (the last o rooms), interpolate from the longest available context down to uniform:

$$\begin{aligned} P_o(x \mid \text{ctx}) &= w \cdot f_o(x \mid \text{ctx}) + (1 - w) \cdot P_{\{o-1\}}(x \mid \text{ctx}[1:]) \\ w &= C(\text{ctx}) / (C(\text{ctx}) + \alpha) \end{aligned}$$

where f_o is the empirical frequency, $C(\text{ctx})$ the times that context was observed, and α a smoothing constant ($\alpha = 1.6, k = 4$). Long contexts dominate once well-supported and decay gracefully when not — Witten-Bell style interpolation, following Begleiter et al. (2004).

Properties that matter for this application:

- **Learns online from the first move** — no pre-training corpus, solving cold start.
- **Persists across runs** — the model carries between deaths; the player's notebook does too.
- **Inspectable** — every prediction decomposes into "context $\rightarrow$ outcome, n observations", which is what makes the death readout and overlay possible.
- **Cheap** — hash-map counter updates, negligible per-frame cost.

3.3 Separating belief from behaviour

The **Predictor** outputs beliefs. The **Director** decides what to do with them, scoring each counter-move by *expected disruption* $\times$ *confidence*:

COUNTER-MOVE	TRIGGER
Intercept	High confidence on a specific room — move to meet the player there
Rewire / rotate	High confidence on a transit route — change it before they reach it
Poison a codebook	A recorded lookup entry the player is predicted to need again
Seal a shortcut	An edge with high traversal frequency

This separation means the model can be upgraded (VOMM $\rightarrow$ GRU $\rightarrow$ STRelay-style architecture) without re-tuning game feel, and game feel can be tuned without touching the model.

3.4 Fairness constraints

An adaptive antagonist must remain defeatable, or the result is unfairness rather than fear (§2.5). Three hard invariants:

1. **No mutation is ever applied to anything currently visible to the player.**
2. **Confidence gating** — below a confidence threshold the antagonist hunts conventionally instead of intercepting. It only exploits what it actually knows.
3. **Solvability invariant** — every topology mutation is checked for graph connectivity and objective reachability, and reverted if it would soft-lock the run. Both mutation operators revert any edit that fails this check, and the property is asserted over 25 generated facilities in the invariant suite (§5.9).

Notably, **no signal is emitted when sabotage occurs**. The player's hand-kept notebook is the only means of detecting it. Sabotage nonetheless lands almost exclusively on recorded information — not as a concession, but because poisoning an unrecorded codebook disrupts nothing, so a utility-maximising Director never spends a move on it. The fairness is emergent, not announced.

3.6 Partial observation: sensing through door events

The antagonist originally observed the player's position perfectly. That turned out to be the reason waiting was not a defence (§5.7), so this cycle replaced it.

Doors carry motion detectors, and a fraction of them are live — some the facility never wired, some the player has spent a dangerous task to silence. The antagonist observes one thing per window: which door, if any, tripped. It maintains a belief over rooms and updates it with a forward filter:

```
predict    b'(c) = sum_r b(r) * T(r -> c)
update     zero any state inconsistent with the evidence, renormalise
```

EVIDENCE	WHAT IT RULES OUT	EFFECT
Door (a,b) tripped	Everything except a crossing of that door	Collapses to near-certainty
Silence	Any crossing of a wired door	Keeps staying put and unwired crossings
Co-location	Everything else	Collapses to certainty

Silence is the interesting case: it is ambiguous between "they held position" and "they used a door I cannot hear", and that ambiguity is what gives the player something to exploit. One consequence is worth stating plainly, because it is counter-intuitive: **if every door is wired, silence proves the player stayed**, so full coverage is still perfect tracking. Partial observation only exists because coverage is incomplete.

3.5 Evaluation protocol

Human playtesting does not scale to a solo project, so evaluation is **headless and bot-driven**. Three scripted players stand in for humans:

- **HabitualBot** — runs a fixed objective circuit (models a player who has learned the level)
- **ExplorerBot** — systematic least-recently-visited sweep (models a first-time player)
- **EvasiveBot** — pursues the same objectives but declines 35% of windows on purpose, to desynchronise from anything modelling it (**the counter-play**)
- **RandomBot** — uniform over doors *and staying* (**the control**: no habit at all)


All accuracy is measured **prequentially** — every prediction is made before its outcome is observed, so there is no train/test leakage by construction.

The control is the most important experiment. Against RandomBot the model must *fail* to beat baseline. A model that appears to predict a random walk is measuring something it should not have access to.

4. Module Description and System Design

4.1 Architecture

Eight modules across three layers. The simulation core has no rendering dependency and is unit-testable in isolation.



4.2 Module responsibilities

MODULE	RESPONSIBILITY	DEPENDS ON
Facility	Room graph; generation; <code>rewire</code> , <code>seal_door</code> ; connectivity and reachability queries. Pure data — the 3D level is built <i>from</i> it.	—
Cipher	Per-room lookup tables, inbound bit sequences, the 9-cell outbound register, transmit validation. Room-agnostic; operates on IDs.	—
Sabotage	The single choke point for all world mutation. Every lie the game tells is logged here and is replayable.	Facility, Cipher
Telemetry	Records room dwell, edge traversals, route repeats, notebook-open events, register-verify duration, flee direction.	—
Predictor	VOMM; <code>predict</code> / <code>observe</code> / <code>top_habit</code> / <code>confidence</code> . Persists across runs.	Telemetry
Director	Scores and selects counter-moves from beliefs. Confidence-gated.	Predictor
Entity	Perception, navigation, patrol/hunt/ambush states. Takes orders; contains no learning .	Director, Facility
Notebook	In-world vector editor (shapes, arrows, freehand, text, frames) and its serialisation. Reads no game state.	—

Two design decisions worth defending:

Sabotage as sole mutator — when debugging "did the AI cheat, or did I misread my own notes?", a single logged mutation path is the difference between a tractable bug and an untraceable one.

Notebook outside the data loop — it is the only channel the Director cannot touch, which is what makes it the player's counter-move rather than decoration.

4.3 The player-facing mechanic

The escape condition is informational, not spatial: transmit an uncorrupted distress signal. Each room carries its **own** lookup table mapping keywords (`HELP US` , `WARNING` , `MAYDAY` , `ALERT` , `CODE RED`) to bit sequences. Because tables legitimately differ per room, cross-referencing two rooms proves nothing — the player's hand-kept notebook is the sole ground truth.

The notebook is **trustworthy but not accurate**: nothing can edit the player's pages, but they record only what was true when written. When the Director poisons a codebook or rotates a room, the notebook does not become corrupted — it becomes **stale**. The player holds a perfect record of a facility that no longer exists.

4.4 Implementation status

SYSTEM	STATUS	WHERE IT RUNS
Facility — room graph, mutation, connectivity invariants	Complete	<code>facility.py</code> , 4 tests
Predictor — variable-order Markov, online, persists	Complete	<code>predictors.py</code>
Director — counter-move selection, confidence-gated	Complete	<code>director.py</code>
Sabotage — sole mutator of world state, logged	Complete	<code>director.py</code>
Telemetry — routes, dwell, edge traversals	Complete	<code>director.py</code>
Entity — perception, pursuit, interception	Complete	<code>director.py</code>
Sensors and belief filter — partial observation	Complete (new)	<code>sensors.py</code>
Notebook — freehand editor, strokes, undo, pages	In the game	<code>parity-game.html</code>
Cipher — codebooks, register, transmit validation	In the game	<code>parity-game.html</code>
Door clock — lockdown / window cycle	In the game	<code>parity-game.html</code>
First-person build — walkable facility with the AI in it	Playable (new)	<code>parity-3d.html</code>
Godot port — engine build, lighting, animation	Not started	—
Audio, content, progression	Not started	—


The complete core loop is playable in first person: move on a shared door clock, read codebooks, write pages, stage and send a byte, and evade an antagonist that models you. What remains is the engine port, audio and content — presentation rather than mechanism.

5. Results

Ten rooms, 1200–1500 windows × 12 seeds, all seeded and reproducible. Evaluation is prequential: every prediction is made before its outcome is observed, so there is no train/test leakage by construction. Training-quality figures (§5.2) use 12 facilities that were never used to choose anything.

5.1 Next-move prediction

At each window the player takes a door **or declines to move**, so the candidate set is $N(v) \cup \{v\}$. Dwell is predicted, not assumed.




Top-1 accuracy by player type and model

PLAYER	NO MODEL	ORDER-1 MARKOV	VOMM
Habitual	39.1%	60.0%	83.6%
Explorer	31.2%	53.0%	69.6%
Evasive (<i>waits to break pattern</i>)	25.4%	44.3%	53.4%
Random (<i>control</i>)	26.9%	25.9%	26.7%

The control is the row to read first. Against a random player the model performs identically to no model at all, which is what a correctly implemented predictor must do. The gains against structured players are therefore learned habit, not leakage.

5.2 How well is the model actually trained?

A predictor is only as good as the ceiling allows, so the ceiling was measured. An oracle that can see the player's own internal state — which objective they are walking to, whether they are mid-page — reaches **91.8%**. The remainder is the player's own randomness and is unpredictable by anything.



Accuracy against the measured ceiling


PREDICTOR	HELD-OUT ACCURACY
No model	48.2%
Order-1 Markov	57.6%
VOMM	85.2% (settled: 89.0%)
Oracle ceiling	91.8%

That is **97% of the achievable signal**. Against a perfectly consistent player the model reaches **100%**, which settles whether it is capable of perfect prediction: it is, and the shortfall is the player's deviation rate.

Three attempts to close the gap, all negative and all retained as results: a grid search over k and α on held-out facilities (no gain over the hand-picked values); tagging the context with the last objective room (+0.1 points); and `GoalVOMM`, which infers the player's current objective from movement direction and scored *worse*. Every mixture weight tried converged to the baseline as the weight went to zero — the tell that the goal signal carries no information the sequence model does not already hold. The variable-order context already encodes the objective cycle.

5.3 Learning speed


Rolling accuracy rises steeply inside the first minute of play, satisfying the cold-start requirement (§1.3) empirically.



Rolling accuracy against observed moves

5.4 Does prediction make a better antagonist?

Embodied pursuit: the entity occupies a room and moves one room per window, so it must intercept rather than teleport.




Captures per 1000 moves by entity behaviour

ENTITY BEHAVIOUR	CAPTURES / 1000 MOVES
Random walk (<i>control</i>)	95
Camp the busiest junction	301
Order-1 Markov director	342
VOMM, one-step targeting	350
VOMM interceptor	362

Hub-camping is included deliberately: a model that cannot beat "sit in the busiest room" has not earned its complexity. Note honestly that the model-driven directors sit within noise of one another — the clear separation is at prediction (§5.1), not at capture rate.

5.5 Does the advantage hold at scale?




Captures against facility size

ROOMS	RANDOM	CAMP THE HUB	ORDER-1	VOMM INTERCEPTOR	ADVANTAGE
10	98	295	332	360	1.22×
16	61	248	290	309	1.25×
24	38	159	213	238	1.49×
32	33	140	189	196	1.40×

Hub-camping is only competitive on a small graph with one dominant junction. As the facility grows it decays fastest and the interceptor's advantage widens.

5.6 Do the world-editing counter-moves matter?


The entity is removed here, isolating topology mutation from hunting.



Navigation cost with and without world editing

Sealing and rewiring ahead of the player raises the cost of completing an objective by **63%** (16.4 → 26.8 moves per lap). Every mutation passed a connectivity and reachability check, so no run was ever soft-locked.

5.7 Partial observation: what the sensors buy (new)




Localisation against sensor coverage

DOORS WIRED	UNIFORM PRIOR	LEARNED PRIOR
100%	99.9%	99.9%
80%	87.2%	86.1%
60%	67.7%	63.0%
40%	54.0%	50.8%
20%	34.8%	36.7%

A fully wired facility leaks everything, for the reason given in §3.6: silence proves the player stayed. Localisation falls roughly linearly as coverage drops.

5.8 What the player can actually do about it (*new*)

Measured per **completed objective**, not per window — a player caught less often only because they achieved less has not defended themselves. All strategies pay the same notebook cost, and the two sensor strategies use the same navigation as their baseline, so the comparison isolates the strategy rather than the movement code.




Captures per objective by player strategy

DOORS WIRED	BASELINE	WAITING	ROUTING AROUND	DISABLING 3 SENSORS
100%	1.04	6.37 (+439%)	1.04 (+0%)	1.99 (+91%)
60%	2.46	4.62 (+80%)	2.33 (-5%)	1.94 (-21%)
30%	2.23	3.55 (+51%)	2.18 (-2%)	2.15 (-4%)

Three findings. **Waiting is catastrophic**, not merely useless. **Routing around sensors barely registers** — the detour costs more exposure than the concealment saves. **Killing sensors on the route you already use works, but only at moderate coverage**: at full coverage nine windows stood still in the open buys too little when every other door still reports you. The counter-play has a sweet spot, which is a design result as much as a measurement.

5.9 Knowing *when* the player stops to write (*the open problem*)



Recall on the dwell event

Once the player writes realistically — a codebook copied once, not re-copied on every pass — dwell becomes rare (4% of windows) and accuracy stops meaning anything: "they always move" scores 96% while learning nothing. On **recall**, the honest measure for a rare event, the model catches **17%** of the moments the player stops, at 26% precision, against a majority-class baseline that catches none by construction.

Four attempts to improve it — including an explicit first-visit feature — all made it worse. Each split the context and lost more to data sparsity than it gained in signal. This is the clearest open problem in the project.

5.10 Invariant test suite

Nine invariants asserted over 25 generated facilities and thousands of simulated moves
(`model/test_invariants.py`, all passing):

INVARIANT	GUARDS AGAINST
Facility connected on construction	Unplayable generated levels
<code>seal_door</code> never disconnects	The antagonist cutting the map in half
<code>rewire</code> never disconnects	The same, via door relocation
Uplink reachable from every room after mutation	Soft-locked, unwinnable runs
Bots move to an adjacent room or stay put	Silent teleportation corrupting telemetry
Predictor output is a proper distribution	Malformed probabilities
VOMM approximates uniform on a random player	Leakage in the evaluation
Director never mutates a visible room	Fairness rule 1
Director never edits below its confidence gate	Fairness rule 2

5.11 The playable build



The first-person build, with the antagonist's reasoning on screen

`demo/parity-3d.html` is the game: eight hexagonal cells generated as a room graph, first-person movement gated by the door clock, three codebooks to record and an uplink to reach, and an antagonist that hunts on the belief filter of §3.6 while narrating its own reasoning on the HUD — what tripped, how sure it is, what it expects next, and whether it is intercepting or sweeping.

An automated playtest drives the same code a player does and completes the game: three codebooks recorded, uplink reached, signal away, having been caught once and recovered, with zero frames stuck on geometry across 6000 frames.

`demo/parity-game.html` is the same antagonist in a systems view, with its belief drawn as heat on the rooms, wired and dead doors distinguished, and a sensor the player can spend three windows to cut.

6. Roadmap

MILESTONE	DELIVERABLE	STATE
M0	Headless simulation core, telemetry, bot harness	Complete
M1	Predictor, baselines, controls, offline metrics	Complete
M2	Director and counter-moves, validated headless	Complete
M6	Door motion-sensors and partial observation	Complete this cycle
M4	Notebook, cipher and door clock, in the playable build	Complete
M3b	First-person WebGL build with the AI inside it	Playable
M3	Godot port — engine tooling, lighting, animation	Next
M5	Human playtest; replace scripted players with real traces	Next
M7	Polygonal room tiles, side count as the difficulty setting	Planned
M8	Outbound register with pre-send tampering, in 3D	Planned

M6 was taken ahead of the engine port because the evidence pointed at it: the negative result in §5.7's predecessor identified perfect observation as the reason waiting was not a defence, and fixing the observation model was worth more than better rendering.

The clearest open problems are **predicting when the player stops to write** (§5.9) and **replacing the simulated player with human traces** (M5). Every behavioural parameter in this report — the deviation rate, the writing model — is currently an assumption, and the playtest is what turns them into measurements.

7. References

1. Lin, Z. et al. (2026). *Unlocking Open-Player-Modeling-enhanced Game-Based Learning: The Open Player Socially Analytical Intelligence Architecture*. arXiv:2603.26915
2. Dehpanah, A. et al. (2021). *Player Modeling using Behavioral Signals in Competitive Online Games*. arXiv:2112.04379
3. Carlier, S. et al. (2024). *Personalization in Serious Games and Gamification for Healthcare: A Three-Tiered Review of Models, Methods and Opportunities*. arXiv:2411.18500
4. Halina, E. et al. (2026). *Representing and Generating Levels Over Time through Playtrace Reconstructive Partitioning*. arXiv:2607.12097
5. Bazzaz, M. (2026). *Player Perceptions of Generative AI in Games: A Steam Review Analysis*. arXiv:2608.11539
6. Romeo, C. et al. (2025). *NTRL: Encounter Generation via Reinforcement Learning for Dynamic Difficulty Adjustment in Dungeons and Dragons*. arXiv:2506.19530
7. Lopes, P. et al. (2025). *Closing the Loop in Affect-Driven Game Adaptation: A Systematic Review*. arXiv:2505.01351
8. Cafri, N. et al. (2025). *Dynamic Difficulty Adjustment With Brain Waves as a Tool for Optimizing Engagement*. arXiv:2504.13965

9. Fuchs, R. et al. (2024). *Personalized Dynamic Difficulty Adjustment — Imitation Learning Meets Reinforcement Learning*. arXiv:2408.06818
10. McConnell, M. et al. (2025). *From Frustration to Fun: An Adaptive Problem-Solving Puzzle Game Powered by Genetic Algorithm*. arXiv:2509.23796
11. Jarra Montoya, O. et al. (2026). *Atari Games Challenge: A Pilot Study on Multimodal Player Experience Assessment*. arXiv:2605.27261
12. Tripathi, P. et al. (2026). *AI-Enabled Serious Games: Integrating Intelligence and Adaptivity in Training Systems*. arXiv:2605.21962
13. Fiuza Vieira, I. et al. (2026). *A Novel Procedural Generation for Level Design of Mansions and Dungeons*. arXiv:2606.03857
14. Özkan, M. B. (2025). *Procedural Game Level Design with Deep Reinforcement Learning*. arXiv:2510.15120
15. Earle, S. et al. (2025). *Video Game Level Design as a Multi-Agent Reinforcement Learning Problem*. arXiv:2510.04862
16. Begleiter, R., El-Yaniv, R. & Yona, G. (2004). *On Prediction Using Variable Order Markov Models*. Journal of Artificial Intelligence Research, 22, 385–421. arXiv:1107.0051
17. Wu, Y. et al. (2025). *Beyond Regularity: Modeling Chaotic Mobility Patterns for Next Location Prediction*. arXiv:2509.11713
18. Deng, B. et al. (2025). *STRelay: A Universal Spatio-Temporal Relaying Framework for Location Prediction over Human Trajectory Data*. arXiv:2508.16620
19. Liu, S. et al. (2025). *Mixture-of-Experts for Personalized and Semantic-Aware Next Location Prediction*. arXiv:2505.24597
20. Soni, A. et al. (2025). *Can We Predict Your Next Move Without Breaking Your Privacy?* arXiv:2507.08843
21. Zhang, H. et al. (2024). *VRMN-bD: A Multi-modal Natural Behavior Dataset of Immersive Human Fear Responses in VR Stand-up Interactive Games*. arXiv:2401.12133
22. Zhang, H. et al. (2023). *Decoding Fear: Exploring User Experiences in Virtual Reality Horror Games*. arXiv:2312.15582